

```
#Name: Asita Ganatra  
#PRN: 22070521003  
#Dataset: Stock Market Dataset
```

```
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
import numpy as np  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression  
from sklearn import metrics  
from PIL import Image  
import os  
  
# Define the file name  
file_name = 'Stock_Market_Dataset.xlsx - Sheet1.csv'  
  
# Load the dataset  
try:  
    df = pd.read_csv(file_name)  
    print("Data loaded successfully.")  
    print(df.head())  
except FileNotFoundError:  
    print(f"Error: File not found at '{file_name}'")
```

```
Error: File not found at 'Stock_Market_Dataset.xlsx - Sheet1.csv'
```

```
# Clean column names  
df.rename(columns={  
    'Close*': 'Close',  
    'Adj Close**': 'Adj_Close'  
}, inplace=True)  
  
# Convert 'Date' column to datetime  
df['Date'] = pd.to_datetime(df['Date'], format='%b %d, %Y')  
  
# Sort by Date  
df.sort_values('Date', inplace=True)  
  
print("\nData after preprocessing:")  
print(df.info())
```

```
Data after preprocessing:  
<class 'pandas.core.frame.DataFrame'>  
Index: 1258 entries, 1257 to 0  
Data columns (total 7 columns):
```

```
#   Column    Non-Null Count Dtype
#   --   --   --
0   Date      1258 non-null   datetime64[ns]
1   Open       1258 non-null   float64
2   High       1258 non-null   float64
3   Low        1258 non-null   float64
4   Close      1258 non-null   float64
5   Adj_Close  1258 non-null   float64
6   Volume     1258 non-null   int64
dtypes: datetime64[ns](1), float64(5), int64(1)
memory usage: 78.6 KB
None
```

```
print("\nGenerating exploratory plots...")
sns.set_style("whitegrid")

# Plot 1: Time Series of Closing Price
plt.figure(figsize=(12, 6))
plt.plot(df['Date'], df['Close'])
plt.title('Stock Closing Price Over Time')
plt.xlabel('Date')
plt.ylabel('Closing Price')
plt.savefig('time_series_close.png')
plt.clf() # Clear the figure

# Plot 2: Scatter Plot (Open vs. Close)
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='Open', y='Close')
plt.title('Relationship Between Opening and Closing Price')
plt.xlabel('Opening Price')
plt.ylabel('Closing Price')
plt.savefig('scatter_open_close.png')
plt.clf() # Clear the figure

print("Saved 'time_series_close.png' and 'scatter_open_close.png'")
```

```
Generating exploratory plots...
Saved 'time_series_close.png' and 'scatter_open_close.png'.
<Figure size 1200x600 with 0 Axes>
<Figure size 1000x600 with 0 Axes>
```

```
# Define X (independent variable) and Y (dependent variable)
X = df[['Open']] # Needs to be a 2D array for sklearn
Y = df['Close']

# Split data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

print(f"\nTraining set size: {X_train.shape[0]} samples")
print(f"Testing set size: {X_test.shape[0]} samples")
```

```
Training set size: 1006 samples
Testing set size: 252 samples
```

```
# Create the linear regression model
model = LinearRegression()

# Train the model
model.fit(X_train, y_train)

print("\nModel training complete.")
```

Model training complete.

```
# Get model coefficients
intercept = model.intercept_
coefficient = model.coef_[0]

print("\n--- Simple Linear Regression Model Results ---")
print(f"Intercept (b): {intercept}")
print(f"Coefficient (m) for 'Open': {coefficient}")
print(f"The model equation is: Close = {coefficient:.4f} * Open

# Make predictions on the test set
y_pred = model.predict(X_test)

# Calculate and print evaluation metrics
r2 = metrics.r2_score(y_test, y_pred)
mae = metrics.mean_absolute_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

print("\n--- Model Performance Metrics (on Test Set) ---")
print(f"R-squared (R²): {r2:.4f}")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
```

--- Simple Linear Regression Model Results ---
 Intercept (b): 100.27809590911056
 Coefficient (m) for 'Open': 0.9965693435593619
 The model equation is: Close = 0.9966 * Open + 100.2781

--- Model Performance Metrics (on Test Set) ---
 R-squared (R²): 0.9942
 Mean Absolute Error (MAE): 210.75
 Root Mean Squared Error (RMSE): 294.45

Start coding or generate with AI.

```
# Step 4: Choose dependent and independent variables
# Rename 'Close*' to 'Close' for consistency with the task descri
df.rename(columns={'Close*': 'Close'}, inplace=True)
```

```
X = df[['Open']] # Independent variable  
y = df['Close'] # Dependent variable  
  
# Step 5: Split dataset into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Step 6: Create and train the model  
model = LinearRegression()  
model.fit(X_train, y_train)
```


```
LinearRegression (1) (?)  
LinearRegression()
```

```
# Step 7: Make predictions  
y_pred = model.predict(X_test)
```

```
# Step 8: Evaluate the model  
print("\nModel Evaluation Metrics:")  
print("Mean Absolute Error:", mean_absolute_error(y_test, y_pred))  
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))  
print("R² Score:", r2_score(y_test, y_pred))
```

```
Model Evaluation Metrics:  
Mean Absolute Error: 204.87196943684768  
Mean Squared Error: 86846.33420129889  
R² Score: 0.9946453497353263
```

```
# Step 9: Visualize results  
plt.figure(figsize=(8,6))  
plt.scatter(X_test, y_test, color='blue', label='Actual Values')  
plt.plot(X_test, y_pred, color='red', linewidth=2, label='Predictions')  
plt.title('Simple Linear Regression')  
plt.xlabel('Open Price')  
plt.ylabel('Close Price')  
plt.legend()  
plt.show()
```



```
# Step 10: Print coefficients
print("\nRegression Equation: Close = {:.4f} * Open + {:.4f}".fo
```

```
Regression Equation: Close = 0.9966 * Open + 100.2781
```

```
#Residual Plot (to check error distribution)
residuals = y_test - y_pred
plt.figure(figsize=(8,6))
sns.scatterplot(x=y_pred, y=residuals, color='purple')
plt.axhline(y=0, color='black', linestyle='--')
plt.title("Residual Plot")
plt.xlabel("Predicted Values")
plt.ylabel("Residuals (Error)")
plt.show()
```

